

# Programming and Data Analysis for Scientists

*Seminar* – Compiled vs Interpreted code

Prof Stephen Clark



University of  
BRISTOL

SCIF20002

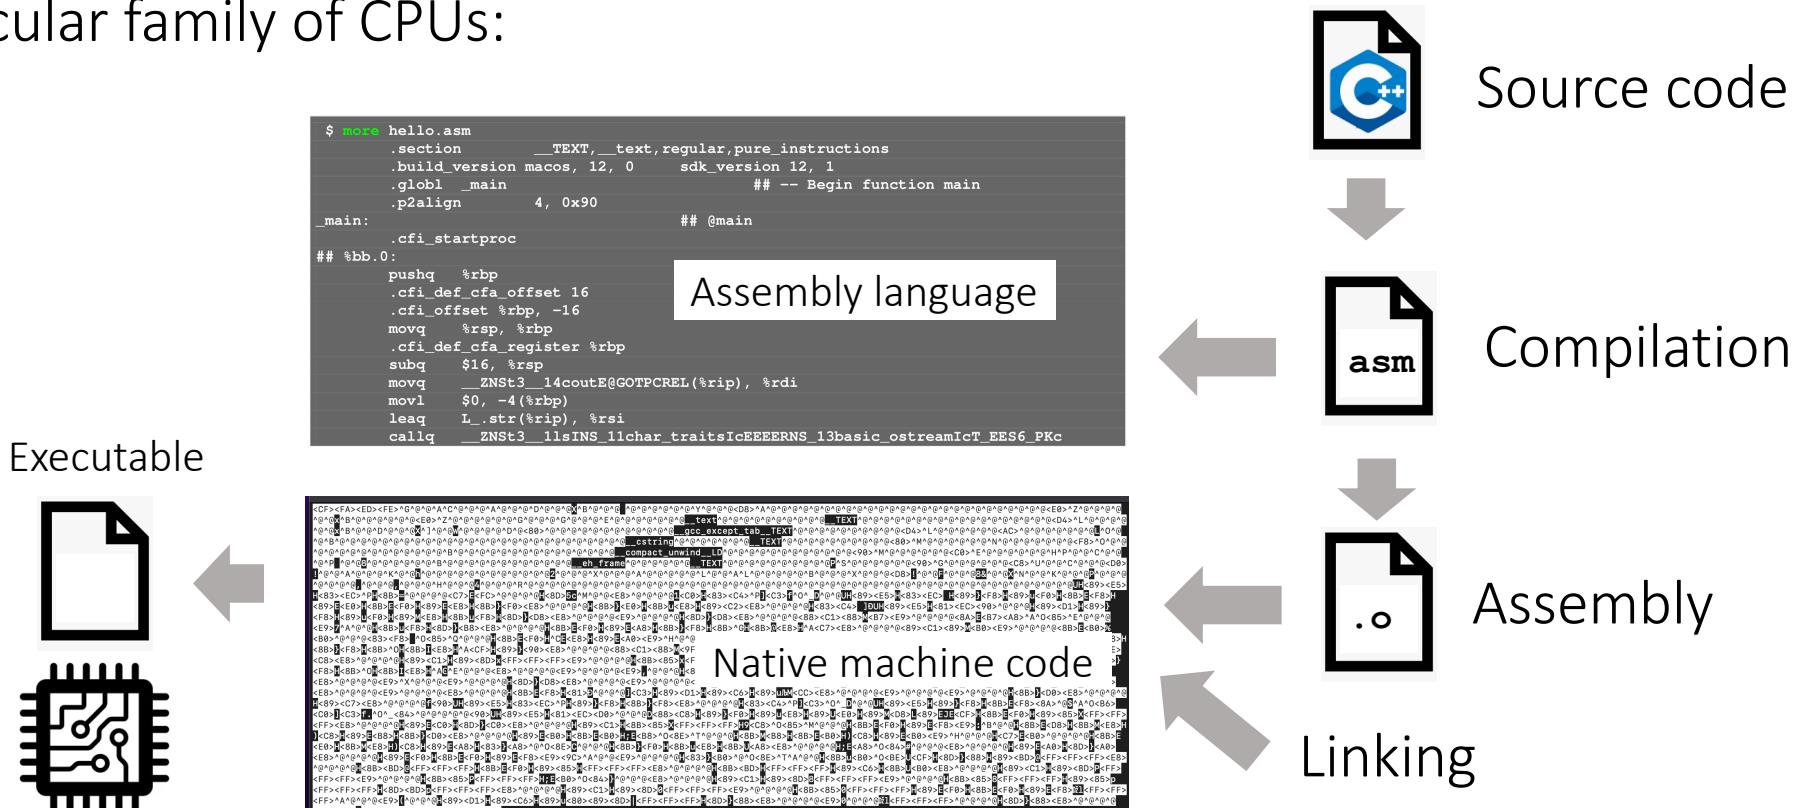
# Compiled vs Interpreted code\*

- We have been learning C/C++ for a couple of weeks now. A key feature of this language is that code is **compiled** making it **compact** and **fast**.
- Python is much more user-friendly than C/C++ making it extremely easy to prototype and explore new ideas and methodologies (with a big ecosystem of libraries).
- **But**, Python is slow due to being **interpreted**. In particular explicit looping is very slow due to overheads created by Python's dynamic variables type.

\***Pedants:** technically whether code is compiled or interpreted is not a property of the *language* itself, but rather of the implementation. You could have a C/C++ interpreter and a Python compiler (although this is difficult to code).

# What is compilation?

Compiling takes source code and converts it into binary machine code native to a particular family of CPUs:



# What is interpretation?

Interpreted source code is usually “compiled” not into binary machine code but into a bytecode – a kind of machine code for a virtual machine (hardware abstracted). The virtual machine is then a compiled executable that “runs” the bytecode instructions.



This form of virtualization approach was pioneered by *Java* and has become very popular because it is hardware agnostic (“write once run anywhere”).

# Simple speed test

Ok, let's do a straightforward speed test between C++ and Python.



Write a Python script (`counter.py`) and a C++ program (`counter.cpp`) which iterates a counter from 0 to 1 billion and then prints the value of the counter to the console at the end.

Run and time them on the command line (Mac terminal shown):

```
% time python counter.py
```

```
% g++ counter.cpp -o counter  
% time ./counter
```

You should get something like this:

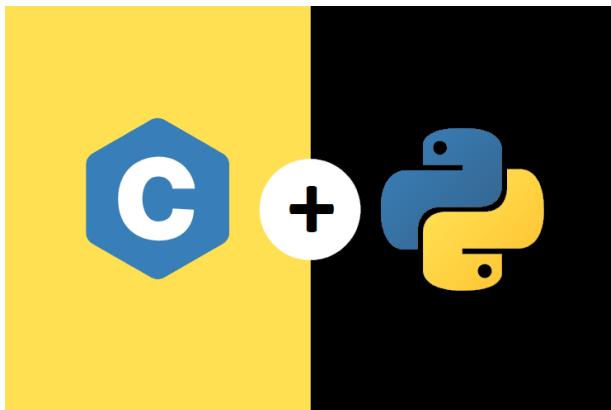
```
1000000000  
python counter.py 43.83s user 0.06s system 99% cpu 43.911 total
```

Python is nearly 100 times slower!

```
1000000000  
.counter 0.52s user 0.00s system 78% cpu 0.672 total
```

# Calling compiled C++ code from Python

If you know which bits of your code are the bottleneck (it is usually clear) then you could write them in C++ (compile it) and then call them from Python.

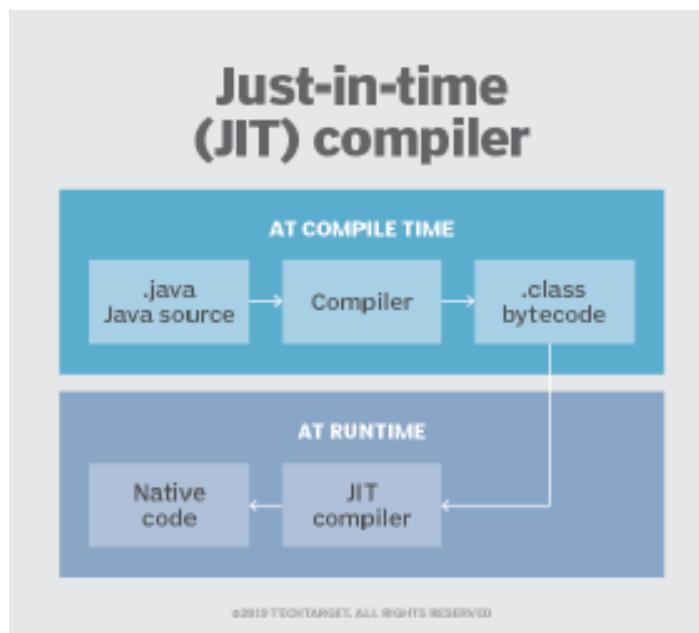


This is a mixed language approach attempting to gain the best of both worlds.

We will have a look at an example using `ctypes` shortly. The issue here is that this kind of binding can be complicated to set up and tricky to debug.

# Just in time compilation

An in-between possibility is so called *just in time* compilation. The virtual machine profiles your code to find the “hot-spots” and then compiles these parts during run-time into native machine code.



For numerical calculations the **Numba** package for Python basically does ctypes for you without any effort or recoding.



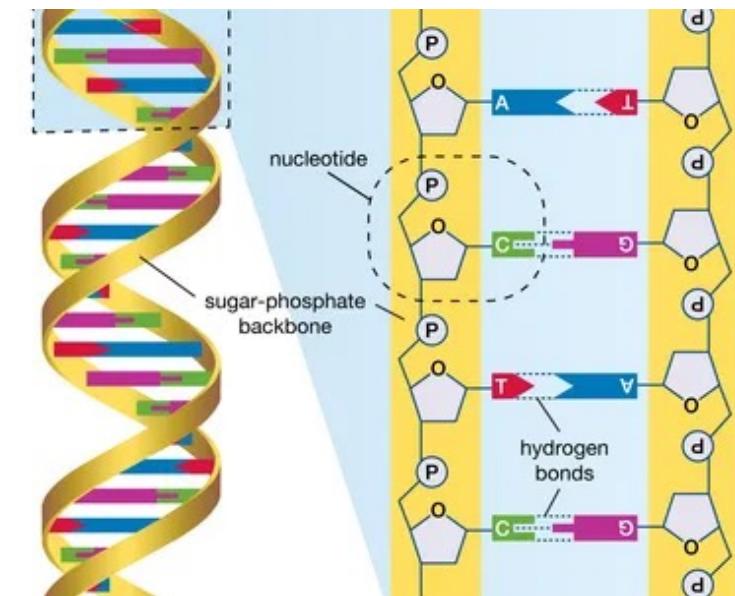
If it works you can get C/C++ similar performance from Python. Let do some more tests ...

# Example 1: Computing DNA K-mers

A DNA is a long chain of units called nucleotides.

There are 4 types of nucleotides labelled by A, C, G, and T. A short portion DNA could look like:

ACTAGGGATCATGAAGATAATGTTGGTGTGTTG



If you choose any K consecutive nucleotides (i.e. letters) from this string, it will be a K-mer. Here are some examples of 4-mers derived from the example.

ACTA, CTAG, TAGG, AGGG, GGGA, etc.

# Example 1: The challenge

Let's generate all possible 13-mers. Mathematically it is a permutation with a replacement problem. Therefore, we have  $4^{13}$  (=67,108,864) possible 13-mers.

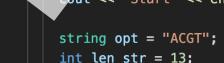
We will use a simple algorithm to generate results in Python and C++ paying particular attention to how long they take to run (they time themselves) ...

```
1 # kmer permutation construction in python:
2
3 from time import time
4
5 # Define a time measuring function:
6 def measure_time(function):
7     def timed(*args, **kwargs):
8         begin = time()
9         result = function(*args, **kwargs)
10        end = time()
11
12        print('Completed in {:.1f} seconds'.format(end - begin))
13        return result
14    return timed
15
16 # Conversion function for nucleotides
17 def convert(c):
18     if (c == 'A'): return 'C'
19     if (c == 'C'): return 'G'
20     if (c == 'G'): return 'T'
21     if (c == 'T'): return 'A'
22
23 # k-mer construction function:
24 @measure_time
25 def construct_kmers(len_str,opt):
26
27     s = ""
28     last = ""
29
30     for i in range(len_str):
31         s += opt[0]
32
33     return s
```



Kmers\_Perm\_Numba.py

```
1 // kmer permutation construction in C/C++
2
3 // g++ --std=c++11 kmers_perm.cpp -o kmers_perm
4 // ./kmers_perm
5
6 #include <iostream>
7 #include <iomanip>
8 #include <string>
9 #include <chrono>
10
11 using namespace std;
12 using namespace std::chrono;
13
14 char convert(char c)
15 {
16     if (c == 'A') return 'C';
17     if (c == 'C') return 'G';
18     if (c == 'G') return 'T';
19     if (c == 'T') return 'A';
20     return ' ';
21 }
22
23 int main()
24 {
25     cout << "Start" << endl;
26
27     string opt = "ACGT";
28     int len_str = 13;
```



Kmers\_Perm.cpp

=> Run the Python code

=> Compile and run the C++ code

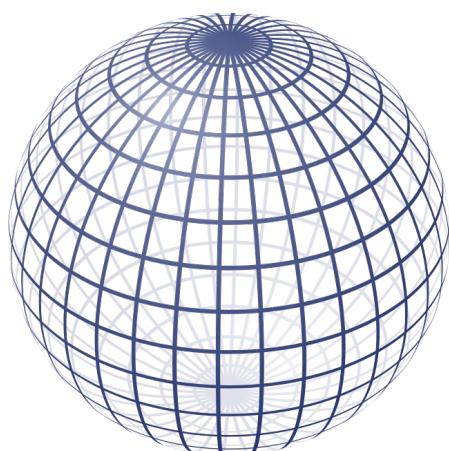
```
g++ --std=c++11 kmers_perm.cpp  
-o kmers_perm
```

```
./kmers_perm
```

How do they compare?

## Example 2: Numerical integration in 3D

Let's try to compute the following 3D integral:



Constant

$$\int_{\Omega} \frac{d\mathbf{k}}{2\pi^2} e^{-\sigma^2 \mathbf{k}^T \boldsymbol{\epsilon} \mathbf{k}}$$

Integrate over a sphere of radius  $G_{\max}$

3D wave-vector

$$\mathbf{k} = \begin{pmatrix} k_x \\ k_y \\ k_z \end{pmatrix}$$

Dielectric permittivity tensor

$$\boldsymbol{\epsilon} = \begin{pmatrix} \epsilon_{xx} & 0 & 0 \\ 0 & \epsilon_{yy} & 0 \\ 0 & 0 & \epsilon_{zz} \end{pmatrix}$$

## Example 2: Numerical integration in 3D

There are clever methods we could use to do this integration, but let's just do brute force discretisation in Cartesian coordinates ...

Use  $N$  points along each axis so the code will involve:

```
loop over  $k_x$ 
    loop over  $k_y$ 
        loop over  $k_z$ 
            if inside sphere then compute
```

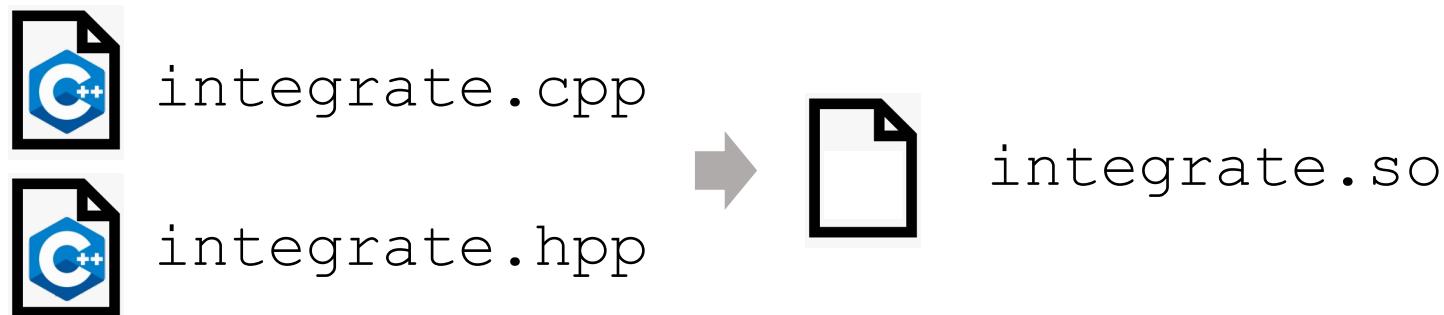
Effort to perform nested loops scales as  $N^3$  ... will become very slow!



## Example 2: Numerical integration in 3D

The code for this example illustrates 3 approaches:

- Fully Python implementation
- Python using `ctypes` binding to a compiled C++ shared library
- C++ code compiling against the C++ shared library

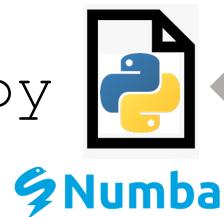


Compile the shared library first as: `g++ -shared integrate.cpp -o integrate.so`

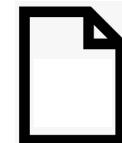
## Example 2: Numerical integration in 3D

Run the python script

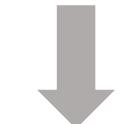
compute\_integral.py



integrate.hpp



integrate.so



and compile the standalone app:

```
g++ --std=c++11 integrate.so compute_integral.cpp -o compute_integral
```

Will need to specify the dynamic link loader path:

```
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:./  
./compute_integral
```



compute\_integral.cpp



compute\_integral

## Example 2: Numerical integration in 3D

How do the approaches compare?

We have 4 configurations to compare:

plain Python	> 1000 s got bored waiting!
Python with Numba	65.7 s
Python using integrate.so via ctypes	8.6 s
plain C++	8.5 s

To get C++ performance ultimately needs C++ compiled code to be used.