

Programming and Data Analysis for Scientists

C++ Workshop 1

The shell and C++

Prof Stephen Clark



University of
BRISTOL

SCIF20002


The shell and C++

The purpose of this workshop is to introduce the C++ programming language. The *learning objectives* are:

- Learn some basic shell command line instructions
- Appreciate the advantages and uses of C++
- Understand what compilation of C++ code means
- Compile, run and post-process data from a simple C++ computation

Using the **shell** terminal

Back in the 1970's a *terminal* was a physical device for logging into a mainframe. These days it refers to a minimal text input and output graphical application emulating this device.

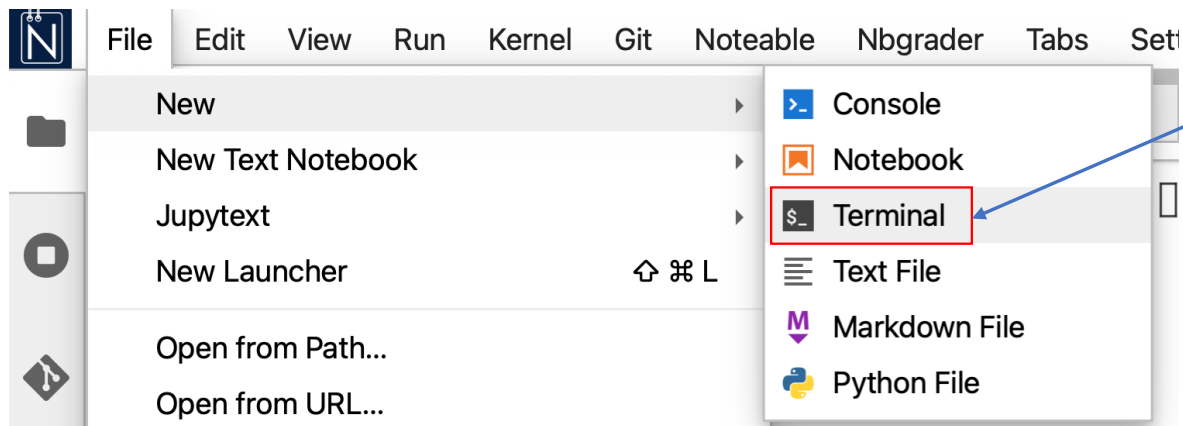


The *shell* is a command-line interpreter program. You type commands into the terminal, which sends them to the shell (like Bash or Zsh) for execution and the return results are displayed back in the terminal window.

While primitive in appearance the shell is extremely powerful. Any interactions with a high-performance computing facility will be via a shell.

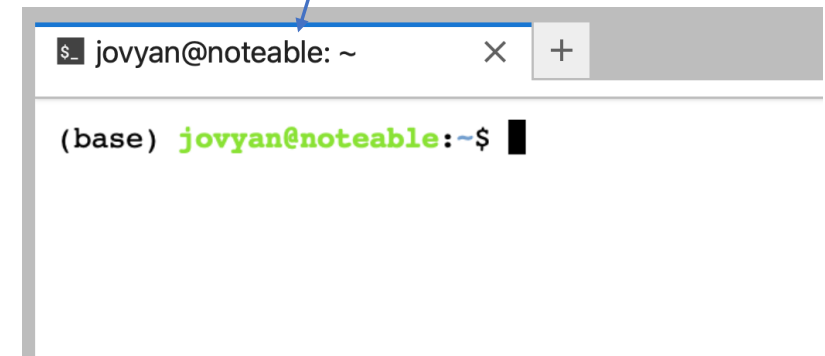
Start a shell on *Noteable*

You are already familiar with the *Noteable* server for running your Python Jupyter lab. The same resource can be used here. Start up Jupyter server ...



Create Terminal session


It appears as a new tab in the Jupyter lab in your browser



We can run shell commands here that will be executed on the Noteable server.

What are C and C++?

C is a procedural programming language developed in 1972. It is a light-weight mid to low-level language developed for systems programming.




C++ is a superset of C first developed in 1985 to enhance it to include object orientated formalism.

C++ uses C syntax along with many new extensions. It continues to be updated and evolve.

Elegant base syntax of C is the *lingua franca* of modern programming. Not only C++, but Java, C#, as well as Python inherited many key features from it.

Why bother to learn C++ in 2025?



Legacy, maturity and ecosystem. It has been around for over 40 years. Unix, Mac OS, Linux, Windows, Google's Android OS are all written and continue to be developed in C++. Most problems have a C++ solution already written.











Since it is compiled it is **fast**! Most new computationally intensive applications like machine learning, scientific computing, virtual machines, device drivers, high-frequency trading, video games ... are all written in C++.





Since it is compiled it can be **small**! For embedded devices powering the IoT memory and processing is limited (due to power consumption and physical size) so their firmware is almost always written in C++.

How popular
are C and C++
currently?

C and C++ still **very** well used.

Jan 2022	Jan 2021	Change	Programming Language		Ratings	Change
1	3	▲	 Python		13.58%	+1.86%
2	1	▼	 C		12.44%	-4.94%
3	2	▼	 Java		10.66%	-1.30%
4	4		 C++		8.29%	+0.73%
5	5		 C#		5.68%	+1.73%
6	6		 Visual Basic		4.74%	+0.90%
7	7		 JavaScript		2.09%	-0.11%
8	11	▲	 Assembly language		1.85%	+0.21%



Source: tiobe.com

Anatomy of a C++ program

Take the customary first program anyone writes `hello.cpp`. We can highlight a number of essential features about C++ code:

It is a “light” language so we often need to import standard libraries to do stuff

All programs start with a `main()` function

```
1 #include <iostream>
2
3 int main()
4 {
5     std::cout << "Hello world!\n"; // The main purpose of this program!
6     return EXIT_SUCCESS; // Return value indicating successful execution
7 }
```

Commands are terminated by ;

It is case-sensitive

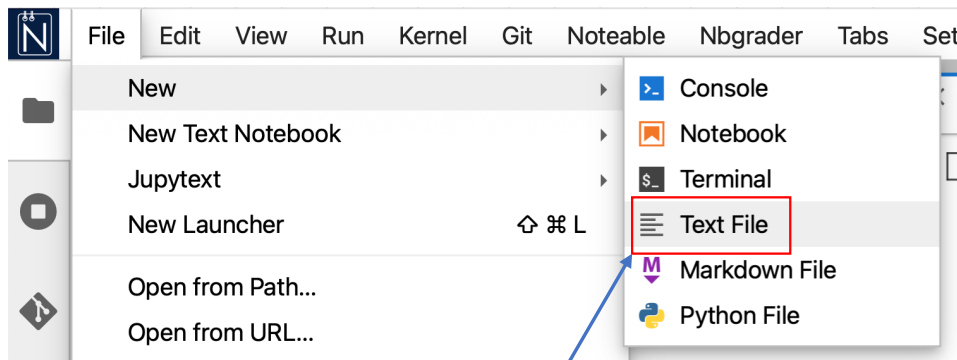
Comments are defined by //

Blocks of code are enclosed by { ... }

Overall, it is a free-form language

Using *Noteable* for C++ code

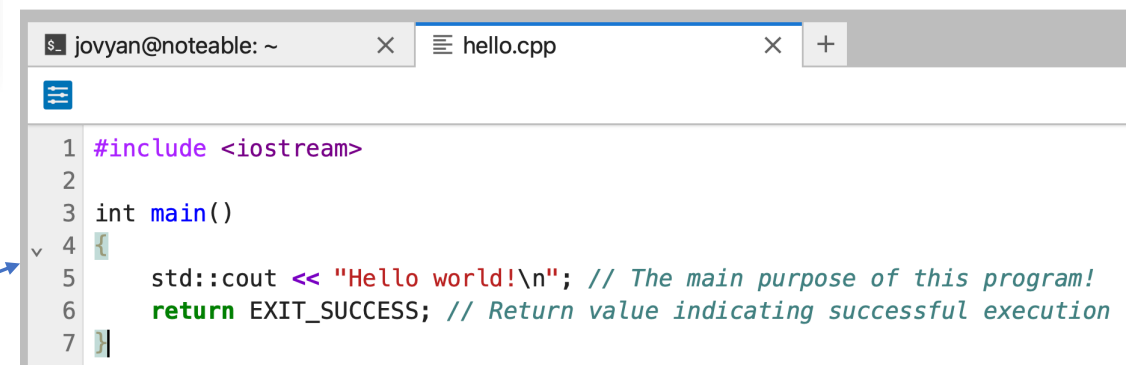
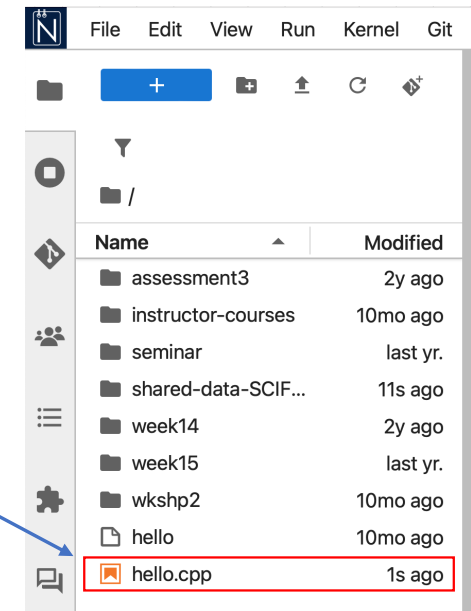
C++ source files are just plain text files with a .cpp extension. We can open or create a file in Noteable and use the built-in code editor ...



Create a new text file opening the editor

New or existing file is opened as a new tab

Open existing file in cloud filesystem



Running C++ code on *Noteable*

When editing a .cpp file C++ language syntax highlighting will automatically switch on. If you created a new text file then rename it <something>.cpp, and then highlighting will turn on:

The image shows the Noteable IDE interface with two windows. The top window, titled 'hello.cpp', displays C++ code with syntax highlighting:

```
1 #include <iostream>
2
3 int main()
4 {
5     std::cout << "Hello world!\n"; // The main purpose of this program!
6     return EXIT_SUCCESS; // Return value indicating successful execution
7 }
```

 A blue arrow labeled 'Language recognition' points to the first line of code. The bottom window, titled 'jovyan@noteable: ~', shows a terminal session. A blue arrow labeled 'Output' points to the terminal output. The terminal shows the following commands and output:

```
(base) jovyan@noteable:~$ ls
assessment3  hello.cpp
hello       instructor-courses  seminar  shared-data-SCIF20002_2023_TB-4  week14  wkshp2
(base) jovyan@noteable:~$ g++ hello.cpp -o hello
(base) jovyan@noteable:~$ ./hello
Hello world!
(base) jovyan@noteable:~$
```

 The commands `g++ hello.cpp -o hello` and `./hello` are highlighted with a red box. A blue arrow labeled 'Compile' points to the first command, and a blue arrow labeled 'Run' points to the second command. The output 'Hello world!' is also highlighted with a red box. To the right of the terminal window, the text 'Switch to shell tab to compile and run in terminal' is displayed.

What is compilation?

```
$ g++ hello.cpp -o hello
```

The process of converting source code into binary machine code. It typically has **4 steps**:

```
$ nm hello.asm
          .section      __TEXT,__text,regular,pure_instructions
          .build_version macos, 12, 0          sdk_version 12, 1
          .globl  _main                          ## -- Begin function main
          .p2align    4, 0x90

          _main:                                          ## @main

          .cfi_startproc

## %bb.0:
          pushq  %rbp
          .cfi_def cfa_offset 16
          .cfi_offset %rbp, -16
          movq   %rsp, %rbp
          .cfi_def cfa_register %rbp
          subq   %16, %rsp
          movq   __ZNSt3_14coutE@GOTPCREL(%rip), %rdi
          movl   $0, -4(%rbp)
          leaq   L_.str(%rip), %rsi
          callq  __ZNSt3_11sINS_11char_traitsIcEEEE8RNSt3__ostreamIcT_EES6_PKc
```

[illegible]

Preprocessing




Compilation



Assembly



Linking



What is HPC?



A “supercomputer” is a specially optimized computer that achieves high performance in floating-point operations (FLOPS) central to scientific calculations.





From the 1990’s onwards supercomputers evolved from being a dedicated single machine to being built from 10,000’s of networked “off-the-shelf” machines allowing for massive parallelization.



Since 2017 all the top 500 supercomputers run a Linux-based OS.

What is HPC?

Currently (2025) the most powerful supercomputers can do in excess of 10^{18} FLOPS distributed over millions of CPU cores.



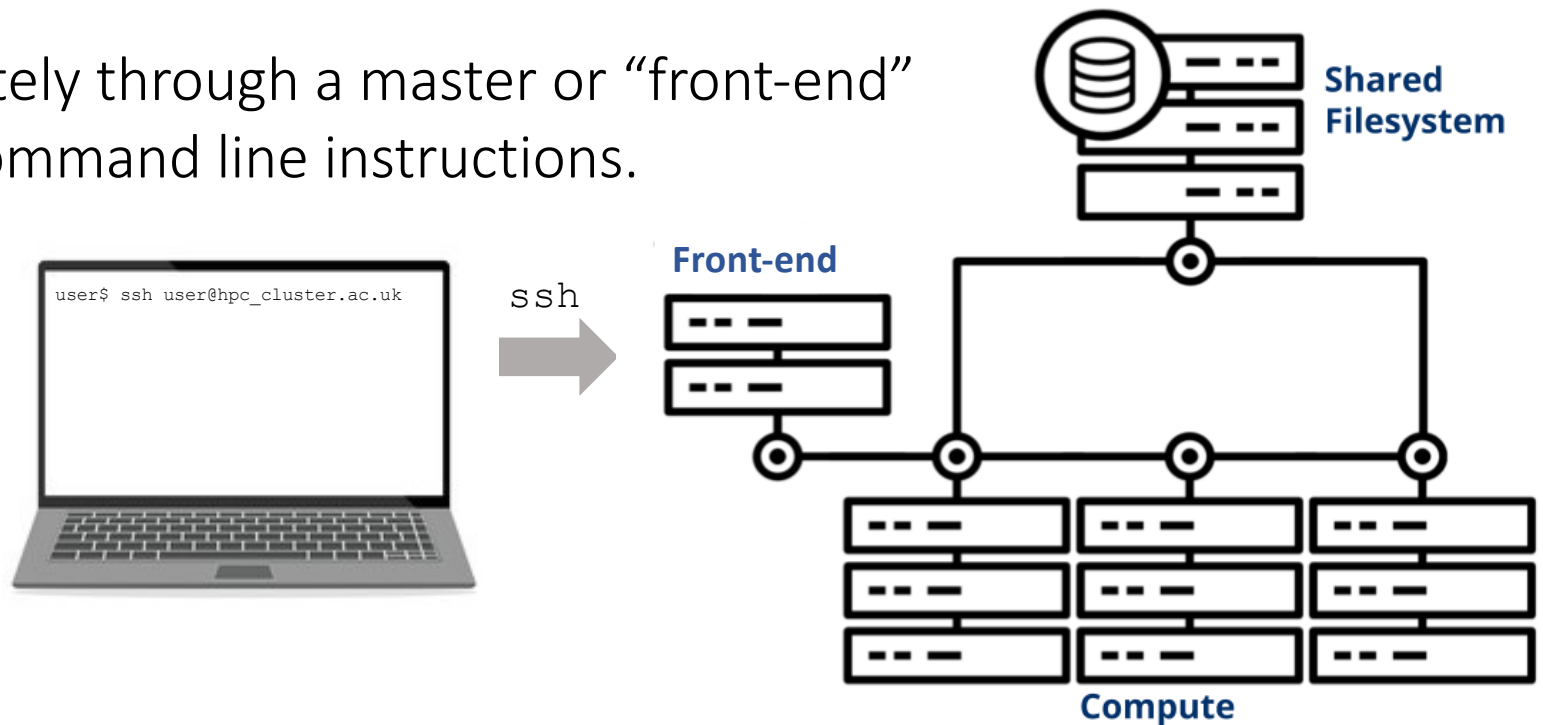
Source: Wikipedia/wiki/Supercomputer

Anatomy of a HPC cluster

Modern HPC cluster comprise many high-powered compute nodes connected via a superfast internal network.


It is accessed remotely through a master or “front-end” node via **ssh** and command line instructions.

Code and data are stored on a shared filesystem.



HPC compute nodes


A “compute node” in a HPC cluster is individually a very high-powered (often dual multicore (12-24) processors) large memory (4-10GB per core) computer packaged compactly as a **rack-mounted** server blade:



Why do we need HPC?

➔ See lecture in Week 15

HPC enables complex scientific calculations that are **impossible** or **infeasibly slow** on a standalone computer. Two cases arise:




- “*Embarrassingly parallelizable problems*”
Naturally divides into many independent tasks whose computation requires no communication between them.

Example: *Monte Carlo sampling*

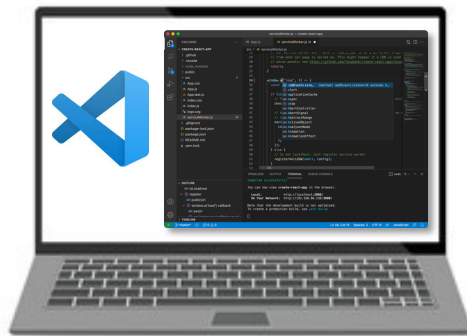
- “*Complex parallelization problems*”
Only divides into many dependent tasks whose computation requires substantial communication between them.

Example: *High-resolution fluid dynamics*




Workflow using HPC

Edit code on your computer



sftp

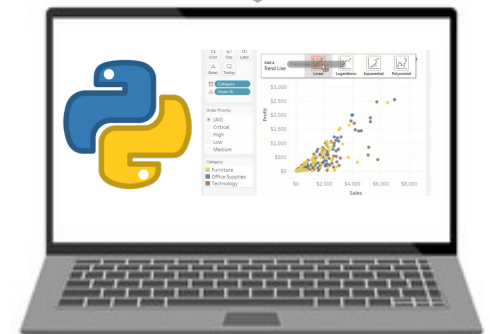



Login



Download data generated

sftp



Compile, test and submit to scheduler

Perform data analysis and visualization on your computer*

A bit heavy handed for simple C++ tests!

Neat example:

You are given C++ code that can compute the well-known Mandelbrot set fractal. The main tasks of this workshop will be to compile it, run it and then post-process the output data to generate a plot similar to these ...

